

✓ Lab 4 Report:

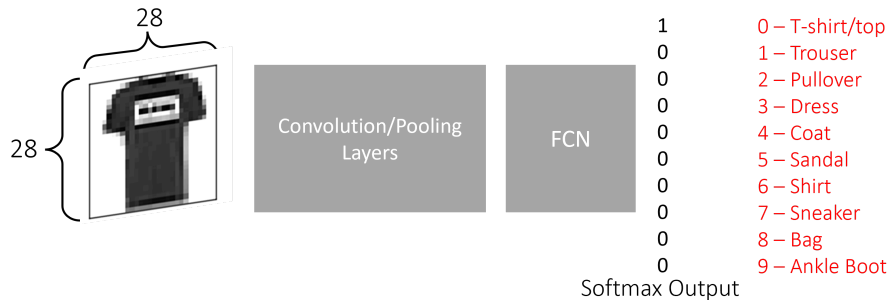
Surpass Human Performance in Fashion MNIST Classification

✓ Name: Travis Hand

```
%matplotlib inline
import matplotlib.pyplot as plt
import torch
import numpy as np
```

```
from IPython.display import Image # For displaying images in colab jupyter cell
```

```
Image('lab4_exercise.png', width = 1000)
```



In this exercise, you will classify fashion item class (28 x 28) using your own **Convolutional Neural Network Architecture**.

Prior to training your neural net, 1) Normalize the dataset using standard scaler and 2) Split the dataset into train/validation/test.

Design your own CNN architecture with your choice of Convolution/Pooling/FCN layers, activation functions, optimization method etc.

Your goal is to achieve a testing accuracy of >89%, with no restrictions on epochs (Human performance: 83.5%).

Demonstrate the performance of your model via plotting the training loss, validation accuracy and printing out the testing accuracy.

After your model has reached the goal, print the accuracy in each class. What is the class that your model performed the worst?

44

✓ Prepare Data

```
# Load Fashion-MNIST Dataset in Numpy
import os
# 10000 training features/targets where each feature is a greyscale image with shape (28, 28)
train_features = np.load('fashion_mnist_train_features.npy')
train_targets = np.load('fashion_mnist_train_targets.npy')

# 1000 testing features/targets
test_features = np.load('fashion_mnist_test_features.npy')
test_targets = np.load('fashion_mnist_test_targets.npy')

# Let's see the shapes of training/testing datasets
print("Training Features Shape: ", train_features.shape)
print("Training Targets Shape: ", train_targets.shape)
print("Testing Features Shape: ", test_features.shape)
print("Testing Targets Shape: ", test_targets.shape)
```

```
↗ Training Features Shape: (10000, 28, 28)
  Training Targets Shape: (10000,)
  Testing Features Shape: (1000, 28, 28)
  Testing Targets Shape: (1000,)
```

```
# Visualizing the first three training features (samples)
```

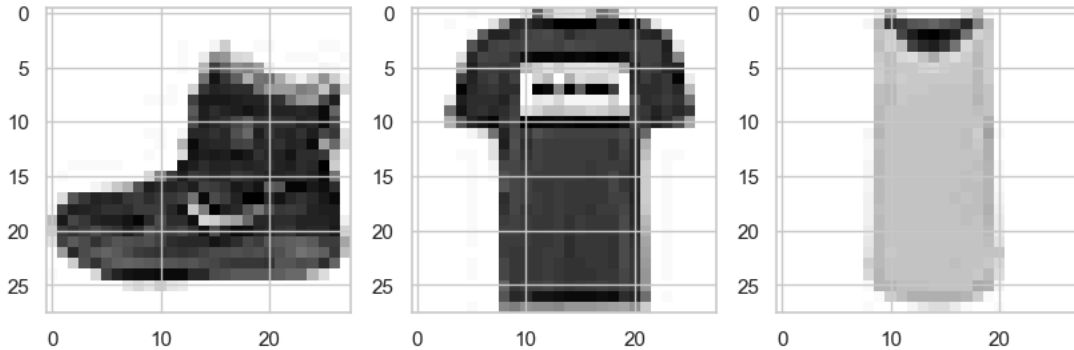
```
plt.figure(figsize = (10, 10))
```

```
plt.subplot(1,3,1)
plt.imshow(train_features[0], cmap = 'Greys')
```

```
plt.subplot(1,3,2)
plt.imshow(train_features[1], cmap = 'Greys')
```

```
plt.subplot(1,3,3)
plt.imshow(train_features[2], cmap = 'Greys')
```

```
<matplotlib.image.AxesImage at 0x165b60920>
```



```
# Reshape features via flattening the images
# This refers to reshape each sample from a 2d array to a 1d array.
# hint: np.reshape() function could be useful here
train_features = train_features.reshape(10000, 784)
test_features = test_features.reshape(1000, 784)
```

```
# Define your scaling function
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
# Scale the dataset according to standard scaling
train_features = scaler.fit_transform(train_features)
test_features = scaler.fit_transform(test_features)
```

```
from sklearn.model_selection import train_test_split
# Take the first 1000 (or randomly select 1000) training features and targets as validation set
# Take the remaining 9000 training features and targets as training set
# Use train_test_split to efficiently split the training targets and tests into validation with 9000:1000 ratio and randomly shuffle
train_features, validation_features, train_targets, validation_targets = train_test_split(train_features, train_targets, test_s
```

```
# Reshape train/validation/test sets to conform to PyTorch's (N, Channels, Height, Width) standard for CNNs
train_features = train_features.reshape(9000, 1, 28, 28)
validation_features = validation_features.reshape(1000, 1, 28, 28)
test_features = test_features.reshape(1000, 1, 28, 28)
```

✓ Define Model

```
# Define your CNN architecture here
```

```
class CNNModel(torch.nn.Module):
```

```
    def __init__(self):
```

```
        super(CNNModel, self).__init__()
```

```
        # Two convolutional layers with 32 and 64 filters respectively
```

```
        # Each convolutional layer is followed by a ReLU activation function
```

```
        self.conv_layer1 = torch.nn.Conv2d(in_channels=1, out_channels=32, kernel_size=3, padding=1) # add stride as well? might
        self.conv_layer2 = torch.nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, padding=1) # maybe increase padding to
```

```
        # Pooling layer to streamline model
```

```
        self.pool = torch.nn.MaxPool2d(2, 2)
```

```

# Fully connected layers for understanding feature relationships
self.fc_layer1 = torch.nn.Linear(64 * 7 * 7, 128)
self.fc_layer2 = torch.nn.Linear(128, 10)

def forward(self, x):

    # Note: If you are using CrossEntropyLoss() do NOT apply softmax to the final output
    # since it's incorporated within the loss function
    x = self.pool(torch.nn.functional.relu(self.conv_layer1(x))) # [N, 32, 14, 14]
    x = self.pool(torch.nn.functional.relu(self.conv_layer2(x))) # [N, 64, 7, 7]
    x = x.view(x.size(0), -1) # Flatten: [N, 64*7*7]
    x = torch.nn.functional.relu(self.fc_layer1(x))
    x = self.fc_layer2(x)
    return x

```

✓ Select Hyperparameters

```

# Fix the random seed so that model performance is reproducible
torch.manual_seed(69)

# Initialize your CNN model

model = CNNModel()

# Define learning rate, epoch and batchsize for mini-batch gradient

learning_rate = 0.001
epochs = 33
batchsize = 50

# Define loss function and optimizer

loss_func = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr = learning_rate, weight_decay=1e-4)

model

↗ CNNModel(
  (conv_layer1): Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv_layer2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc_layer1): Linear(in_features=3136, out_features=128, bias=True)
  (fc_layer2): Linear(in_features=128, out_features=10, bias=True)
)

```

✓ Identify Tracked Values

```

# Placeholders for training loss and validation accuracy during training
# Training loss should be tracked for each iteration (1 iteration -> single forward pass to the network)
# Validation accuracy should be evaluated every 'Epoch' (1 epoch -> full training dataset)
# If using batch gradient, 1 iteration = 1 epoch

train_loss_list = []
validation_accuracy_list = []

```

✓ Train Model

```

import tqdm # Use "for epoch in tqdm.trange(epochs):" to see the progress bar

# Convert the training, validation, testing dataset (NumPy arrays) into torch tensors
# Split your training features/targets into mini-batches if using mini-batch gradient

test_features = torch.tensor(test_features, dtype=torch.float32)
test_targets = torch.tensor(test_targets, dtype=torch.long)
train_features = torch.tensor(train_features, dtype=torch.float32)
train_targets = torch.tensor(train_targets, dtype=torch.long)
validation_features = torch.tensor(validation_features, dtype=torch.float32)
validation_targets = torch.tensor(validation_targets, dtype=torch.long)

```

```

# Create TensorDatasets (feature, target)
train = torch.utils.data.TensorDataset(train_features, train_targets)
validation = torch.utils.data.TensorDataset(validation_features, validation_targets)
test = torch.utils.data.TensorDataset(test_features, test_targets)

# Create data loaders, using batches for more efficient training
train_loader = torch.utils.data.DataLoader(train, batch_size=batchsize, shuffle=True)
val_loader = torch.utils.data.DataLoader(validation, batch_size=batchsize)
test_loader = torch.utils.data.DataLoader(test, batch_size=batchsize)

# select hardware based on availability
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
if torch.cuda.is_available():
    model.cuda()

# Training Loop -----
for epoch in tqdm.trange(epochs):

    model.train()
    train_loss = 0
    correct = 0
    total = 0

    for inputs, targets in train_loader:
        inputs, targets = inputs.to(device), targets.to(device)

        optimizer.zero_grad()
        outputs = model(inputs)

        loss = loss_func(outputs, targets)
        loss.backward()
        optimizer.step()

        train_loss += loss.item() * inputs.size(0)
        predicted = torch.argmax(outputs, 1)
        correct += (predicted == targets).sum().item()
        total += targets.size(0)

    avg_train_loss = train_loss / total
    accuracy = correct / total
    train_loss_list.append(avg_train_loss)

# Compute Validation Accuracy -----

model.eval()
val_loss = 0
val_correct = 0
val_total = 0

with torch.no_grad():
    for val_inputs, val_targets in val_loader:
        val_inputs, val_targets = val_inputs.to(device), val_targets.to(device)

        val_outputs = model(val_inputs)
        loss = loss_func(val_outputs, val_targets)

        val_loss += loss.item() * val_inputs.size(0)
        val_predicted = torch.argmax(val_outputs, dim=1)
        val_correct += (val_predicted == val_targets).sum().item()
        val_total += val_targets.size(0)

    avg_val_loss = val_loss / val_total
    val_accuracy = val_correct / val_total
    validation_accuracy_list.append(val_accuracy)

print(f"Epoch {epoch+1}/{epochs} - "
      f"Train Loss: {avg_train_loss:.4f}, Accuracy: {accuracy:.4f} - "
      f"Val Loss: {avg_val_loss:.4f}, Val Accuracy: {val_accuracy:.4f}")

```

```

/var/folders/nn/yzl4s0_97y791sb31vclsv7w0000gn/T/ipykernel_77070/2139800984.py:6: UserWarning: To copy construct from a tens
test_features = torch.tensor(test_features, dtype=torch.float32)
/var/folders/nn/yzl4s0_97y791sb31vclsv7w0000gn/T/ipykernel_77070/2139800984.py:7: UserWarning: To copy construct from a tens
test_targets = torch.tensor(test_targets, dtype=torch.long)
/var/folders/nn/yzl4s0_97y791sb31vclsv7w0000gn/T/ipykernel_77070/2139800984.py:8: UserWarning: To copy construct from a tens

```

```

train_features = torch.tensor(train_features, dtype=torch.float32)
/var/folders/nn/yzl4s0_97y791sb31vclsv7w0000gn/T/ipykernel_77070/2139800984.py:9: UserWarning: To copy construct from a tens
train_targets = torch.tensor(train_targets, dtype=torch.long)
/var/folders/nn/yzl4s0_97y791sb31vclsv7w0000gn/T/ipykernel_77070/2139800984.py:10: UserWarning: To copy construct from a ten
validation_features = torch.tensor(validation_features, dtype=torch.float32)
/var/folders/nn/yzl4s0_97y791sb31vclsv7w0000gn/T/ipykernel_77070/2139800984.py:11: UserWarning: To copy construct from a ten
validation_targets = torch.tensor(validation_targets, dtype=torch.long)
3%|  | 1/33 [00:03<01:53, 3.55s/it]Epoch 1/33 - Train Loss: 0.6822, Accuracy: 0.7553 - Val Loss: 0.4578, Val Accu
6%|  | 2/33 [00:06<01:45, 3.42s/it]Epoch 2/33 - Train Loss: 0.4120, Accuracy: 0.8542 - Val Loss: 0.4199, Val Accu
9%|  | 3/33 [00:10<01:39, 3.32s/it]Epoch 3/33 - Train Loss: 0.3298, Accuracy: 0.8807 - Val Loss: 0.3747, Val Accu
12%|  | 4/33 [00:13<01:34, 3.28s/it]Epoch 4/33 - Train Loss: 0.2749, Accuracy: 0.8996 - Val Loss: 0.3422, Val Accu
15%|  | 5/33 [00:16<01:30, 3.23s/it]Epoch 5/33 - Train Loss: 0.2406, Accuracy: 0.9108 - Val Loss: 0.3203, Val Accu
18%|  | 6/33 [00:19<01:26, 3.21s/it]Epoch 6/33 - Train Loss: 0.2060, Accuracy: 0.9249 - Val Loss: 0.3434, Val Accu
21%|  | 7/33 [00:22<01:22, 3.19s/it]Epoch 7/33 - Train Loss: 0.1733, Accuracy: 0.9370 - Val Loss: 0.3334, Val Accu
24%|  | 8/33 [00:25<01:19, 3.19s/it]Epoch 8/33 - Train Loss: 0.1489, Accuracy: 0.9451 - Val Loss: 0.3335, Val Accu
27%|  | 9/33 [00:29<01:16, 3.18s/it]Epoch 9/33 - Train Loss: 0.1229, Accuracy: 0.9533 - Val Loss: 0.3451, Val Accu
30%|  | 10/33 [00:32<01:13, 3.19s/it]Epoch 10/33 - Train Loss: 0.1070, Accuracy: 0.9589 - Val Loss: 0.3700, Val Accu
33%|  | 11/33 [00:35<01:10, 3.19s/it]Epoch 11/33 - Train Loss: 0.0921, Accuracy: 0.9664 - Val Loss: 0.3602, Val Accu
36%|  | 12/33 [00:38<01:07, 3.20s/it]Epoch 12/33 - Train Loss: 0.0773, Accuracy: 0.9720 - Val Loss: 0.3700, Val Accu
39%|  | 13/33 [00:41<01:03, 3.17s/it]Epoch 13/33 - Train Loss: 0.0581, Accuracy: 0.9798 - Val Loss: 0.4045, Val Accu
42%|  | 14/33 [00:44<00:59, 3.16s/it]Epoch 14/33 - Train Loss: 0.0475, Accuracy: 0.9840 - Val Loss: 0.4255, Val Accu
45%|  | 15/33 [00:48<00:56, 3.15s/it]Epoch 15/33 - Train Loss: 0.0399, Accuracy: 0.9884 - Val Loss: 0.4013, Val Accu
48%|  | 16/33 [00:51<00:53, 3.15s/it]Epoch 16/33 - Train Loss: 0.0334, Accuracy: 0.9891 - Val Loss: 0.5454, Val Accu
52%|  | 17/33 [00:54<00:50, 3.18s/it]Epoch 17/33 - Train Loss: 0.0437, Accuracy: 0.9841 - Val Loss: 0.4787, Val Accu
55%|  | 18/33 [00:57<00:47, 3.17s/it]Epoch 18/33 - Train Loss: 0.0380, Accuracy: 0.9878 - Val Loss: 0.4929, Val Accu
58%|  | 19/33 [01:00<00:44, 3.17s/it]Epoch 19/33 - Train Loss: 0.0289, Accuracy: 0.9928 - Val Loss: 0.5302, Val Accu
61%|  | 20/33 [01:03<00:41, 3.16s/it]Epoch 20/33 - Train Loss: 0.0245, Accuracy: 0.9932 - Val Loss: 0.5344, Val Accu
64%|  | 21/33 [01:07<00:37, 3.16s/it]Epoch 21/33 - Train Loss: 0.0210, Accuracy: 0.9940 - Val Loss: 0.5587, Val Accu
67%|  | 22/33 [01:10<00:34, 3.16s/it]Epoch 22/33 - Train Loss: 0.0313, Accuracy: 0.9892 - Val Loss: 0.5123, Val Accu
70%|  | 23/33 [01:13<00:31, 3.15s/it]Epoch 23/33 - Train Loss: 0.0157, Accuracy: 0.9958 - Val Loss: 0.5500, Val Accu
73%|  | 24/33 [01:16<00:28, 3.15s/it]Epoch 24/33 - Train Loss: 0.0221, Accuracy: 0.9938 - Val Loss: 0.6120, Val Accu
76%|  | 25/33 [01:19<00:25, 3.15s/it]Epoch 25/33 - Train Loss: 0.0318, Accuracy: 0.9917 - Val Loss: 0.5160, Val Accu
79%|  | 26/33 [01:22<00:22, 3.15s/it]Epoch 26/33 - Train Loss: 0.0320, Accuracy: 0.9890 - Val Loss: 0.5403, Val Accu
82%|  | 27/33 [01:25<00:18, 3.16s/it]Epoch 27/33 - Train Loss: 0.0196, Accuracy: 0.9933 - Val Loss: 0.5422, Val Accu
85%|  | 28/33 [01:29<00:16, 3.20s/it]Epoch 28/33 - Train Loss: 0.0081, Accuracy: 0.9977 - Val Loss: 0.5974, Val Accu
88%|  | 29/33 [01:32<00:13, 3.34s/it]Epoch 29/33 - Train Loss: 0.0027, Accuracy: 0.9998 - Val Loss: 0.6259, Val Accu
91%|  | 30/33 [01:36<00:09, 3.29s/it]Epoch 30/33 - Train Loss: 0.0008, Accuracy: 1.0000 - Val Loss: 0.6379, Val Accu
94%|  | 31/33 [01:40<00:06, 3.47s/it]Epoch 31/33 - Train Loss: 0.0006, Accuracy: 1.0000 - Val Loss: 0.6422, Val Accu
97%|  | 32/33 [01:43<00:03, 3.39s/it]Epoch 32/33 - Train Loss: 0.0005, Accuracy: 1.0000 - Val Loss: 0.6444, Val Accu
100%|  | 33/33 [01:46<00:00, 3.22s/it]Epoch 33/33 - Train Loss: 0.0005, Accuracy: 1.0000 - Val Loss: 0.6458, Val Accu

```

✓ Visualize & Evaluate Model

```

# Seaborn for prettier plot

import seaborn as sns

sns.set(style = 'whitegrid', font_scale = 1)

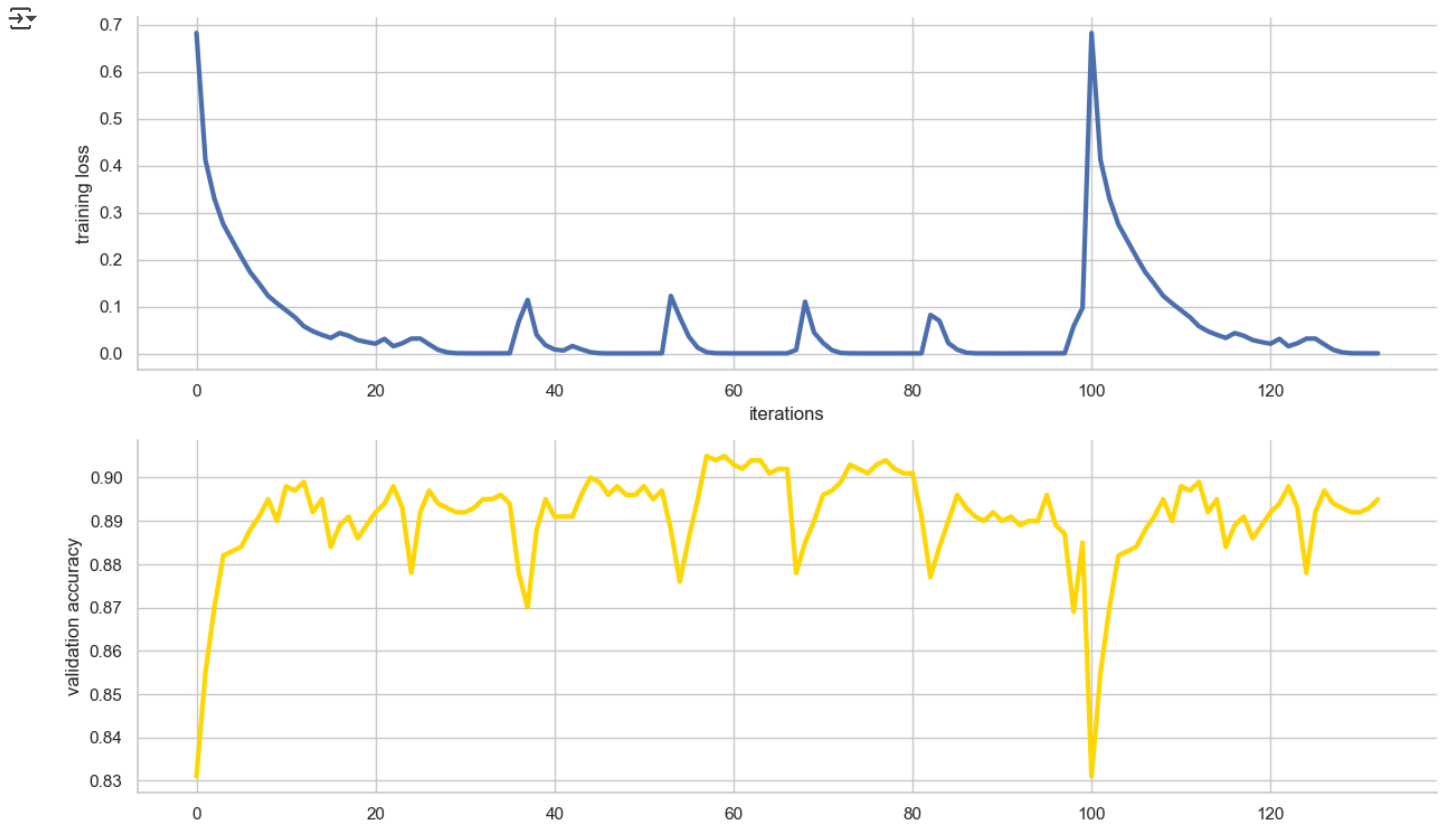
# Visualize training loss

plt.figure(figsize = (15, 9))

plt.subplot(2, 1, 1)
plt.plot(train_loss_list, linewidth = 3)
plt.ylabel("training loss")
plt.xlabel("iterations")
sns.despine()

plt.subplot(2, 1, 2)
plt.plot(validation_accuracy_list, linewidth = 3, color = 'gold')
plt.ylabel("validation accuracy")
sns.despine()

```



```
# Compute the testing accuracy

# Set model to evaluation mode
model.eval()

# Track test accuracy
test_correct = 0
test_total = 0
test_loss = 0.0

with torch.no_grad():
    for inputs, targets in test_loader:
        inputs, targets = inputs.to(device), targets.to(device)

        outputs = model(inputs)
        loss = loss_func(outputs, targets)

        test_loss += loss.item() * inputs.size(0) # sum loss over batch
        predicted = torch.argmax(outputs, 1)
        test_correct += (predicted == targets).sum().item()
        test_total += targets.size(0)

# Final metrics
avg_test_loss = test_loss / test_total
test_accuracy = test_correct / test_total

print(f" Test Loss: {avg_test_loss:.4f}")
print(f" Test Accuracy: {test_accuracy:.4f}")
```

```
Test Loss: 0.7045
Test Accuracy: 0.8960
```

```
# (OPTIONAL) Print the testing accuracy for each fashion class. Your code should produce something that looks like:
# Clever usage of np.where() could be useful here
```

```
# "Accuracy of T-shirt/top: 93.5 %"
```

```

# "Accuracy of Trouser: 89.3 %"
# etc...

# What's the fashion item that your model had the hardest time classifying?

import numpy as np

# Switch to evaluation mode and move model to correct device
model.eval()
all_preds = []
all_targets = []

with torch.no_grad():
    for inputs, targets in test_loader:
        inputs = inputs.to(device)
        outputs = model(inputs)
        preds = torch.argmax(outputs, dim=1)

        all_preds.append(preds.cpu())
        all_targets.append(targets)

# Concatenate predictions and targets into single arrays
all_preds = torch.cat(all_preds).numpy()
all_targets = torch.cat(all_targets).numpy()
clothing = {
    0 : 'T-shirt/top',
    1 : 'Trouser',
    2 : 'Pullover',
    3 : 'Dress',
    4 : 'Coat',
    5 : 'Sandal',
    6 : 'Shirt',
    7 : 'Sneaker',
    8 : 'Bag',
    9 : 'Ankle Boot'
}
least_acc = [1, -1]

# Compute per-class accuracy using np.where
for cls in range(10):
    idx = np.where(all_targets == cls)[0]
    total = len(idx)
    correct = np.sum(all_preds[idx] == all_targets[idx])
    acc = correct / total if total > 0 else 0
    least_acc = least_acc if acc > least_acc[0] else [acc, cls]
    print(f"Class {clothing[cls]}:\t\t Accuracy = {acc:.2%}")

print(f'\nModel had hardest time classifying {clothing[least_acc[1]]}s')

```

```

↩ Class T-shirt/top:          Accuracy = 89.72%
  Class Trouser:              Accuracy = 98.10%
  Class Pullover:             Accuracy = 85.59%
  Class Dress:                Accuracy = 89.25%
  Class Coat:                 Accuracy = 77.39%
  Class Sandal:               Accuracy = 97.70%
  Class Shirt:                Accuracy = 73.20%
  Class Sneaker:              Accuracy = 97.89%
  Class Bag:                  Accuracy = 96.84%
  Class Ankle Boot:           Accuracy = 93.68%

```

Model had hardest time classifying Shirts